

# Dynamic Process Isolation

Black Hat Asia 2022

**Martin Schwarzl**

Graz University of Technology

**Kenton Varda**

Cloudflare Inc.

**Michael Schwarz**

CISPA Helmholtz Center for Information Security

**Pietro Borrello**

Sapienza University of Rome

**Thomas Schuster**

Graz University of Technology

**Andreas Kogler**


Graz University of Technology

**Daniel Gruss**

Graz University of Technology




- **Cloudflare Workers** is one of the top edge-computing systems





- **Cloudflare Workers** is one of the top edge-computing systems
- **Single-process** design



Virtual machine



Isolate model



User code



Process overhead



- **Cloudflare Workers** is one of the top edge-computing systems
- **Single-process** design
- Vulnerable to **Spectre** attacks?



- Cloudflare Workers is one of the top edge-computing systems
- Single-process design
- Vulnerable to Spectre attacks?
- No local timers





- **Cloudflare Workers** is one of the top edge-computing systems
- **Single-process** design
- Vulnerable to **Spectre** attacks?
- No **local timers**
- Number of memory, (sub-)requests, runtime is **limited**



- **Cloudflare Workers** is one of the top edge-computing systems
- **Single-process** design
- Vulnerable to **Spectre** attacks?
- No **local timers**
- Number of memory, (sub-)requests, runtime is **limited**
- If an attack is possible, we need a **low-overhead** solution!

```
maccess(i);
```

```
maccess(i);
```







```
maccess(i);
```


```
maccess(i);
```

*Cache miss*




```
maccess(i);
```

```
maccess(i);
```




```
maccess(i);
```


```
maccess(i);
```





```
maccess(i);
```

```
maccess(i);
```










Shared Memory

ATTACKER


flush  
access





VICTIM


access




















```
if (index < data_size)
{
    y = lut[data[index]*4096]
}
```




`index = 0`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z

```
if (index < 4)
  then lut[data[index]]
  else {
```



Memory


D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	

```
index = 0
```

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z

```
if (index < 4)
  then lut[data[index]]
  else Speculate {}
```




Memory


D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	


`index = 0`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z








`index = 1`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z

```
if (index < 4)
  then lut[data[index]]
  else {}
```





Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	

`index = 1`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z




Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	


`index = 1`


Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z










`index = 1`


Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z



Memory



D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	



`index = 2`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z




Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	


index = 2


Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z



index = 2







`index = 2`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z



Memory



D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	



`index = 3`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z




Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	


index = 3

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z





A



Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	







`index = 3`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z




Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	

`index = 4`


Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z



Memory


D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	




`index = 4`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z




K




Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	









`index = 4`

Lookup Table

	A	B
C	D	E
F	G	H
I	J	K
L	M	N
O	P	Q
R	S	T
U	V	W
X	Y	Z



Memory

D	data[0]
A	data[1]
T	data[2]
A	data[3]
K	
E	
Y	
...	



Application-Level



Application-Level



System-Level



Application-Level



System-Level



Hardware-Level



- Freeze the time during execution



- Freeze the time during execution
- No native code





- Freeze the time during execution
- No native code
- No shared-memory








- Freeze the time during execution
- No native code
- No shared-memory
- No multithreading














- **Amplify** timing either:






- **Amplify** timing either:
  - Encode secret into **multiple** cache lines






- **Amplify** timing either:
  - Encode secret into **multiple** cache lines
  - Loop over gadget  $n$  times




- **Amplify** timing either:


- Encode secret into **multiple** cache lines
- Loop over gadget  $n$  times

```
for (int i = 0; i < N; i++) {  
    // evict A and B  
    // ...  
    if (secret_bit) { access A } // transient  
    else { access B }  
    access A  
}
```







- No **native** code execution




- No **native** code execution
  - Cannot **Flush & Reload** memory



- No **native** code execution
  - Cannot **Flush & Reload** memory
  - Cannot even build **eviction sets**




- No **native** code execution
  - Cannot **Flush & Reload** memory
  - Cannot even build **eviction sets**
  - ✓ Evict the **whole** cache iterating over a huge array




- JavaScript is **optimized** and **deoptimized**







- JavaScript is **optimized** and **deoptimized**
  - Use **assumptions** on variables



- JavaScript is **optimized** and **deoptimized**
  - Use **assumptions** on variables
  - Deoptimize if **invalidated**




- JavaScript is **optimized** and **deoptimized**
  - Use **assumptions** on variables
  - Deoptimize if **invalidated**
  - ✓ Hide the attack behind a **mispredicted** branch




- JavaScript is **optimized** and **deoptimized**
  - Use **assumptions** on variables
  - Deoptimize if **invalidated**
  - ✓ Hide the attack behind a **mispredicted** branch
  - ✓ **Prevent** inlining and further optimizations using huge functions




- JavaScript objects cannot index the **whole** memory



- JavaScript objects cannot index the **whole** memory
- To leak from different workers you need **64-bit addresses**



- JavaScript objects cannot index the **whole** memory
- To leak from different workers you need **64-bit addresses**
- ✓ Use **ArrayBuffers** to leak the whole 64-bit address space



- JavaScript objects cannot index the **whole** memory
- To leak from different workers you need **64-bit addresses**
- ✓ Use **ArrayBuffers** to leak the whole 64-bit address space
- ✓ Create objects of different types and **confuse** branch predictor




$\text{map}_p$
$\text{prototype}_p$
0x1337
...
...
ArrayBuffer*

Object P


```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: new ArrayBuffer(0x1000),  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



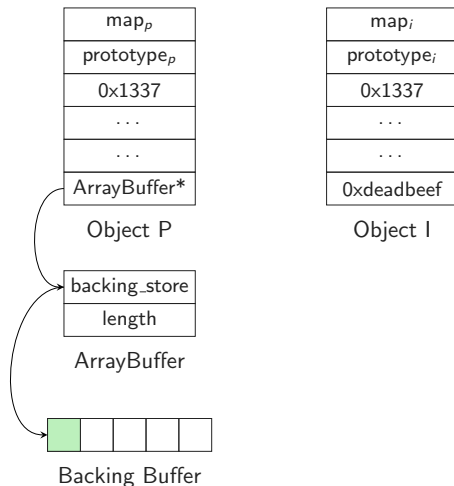
```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: new ArrayBuffer(0x1000),  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



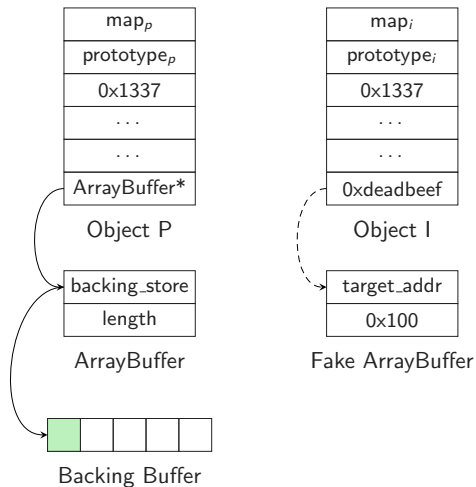
```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: new ArrayBuffer(0x1000),  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```




```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: 0xdeadbeef,  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



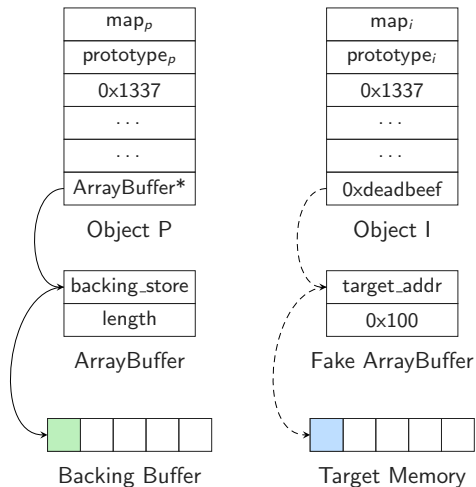
```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: 0xdeadbeef,  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



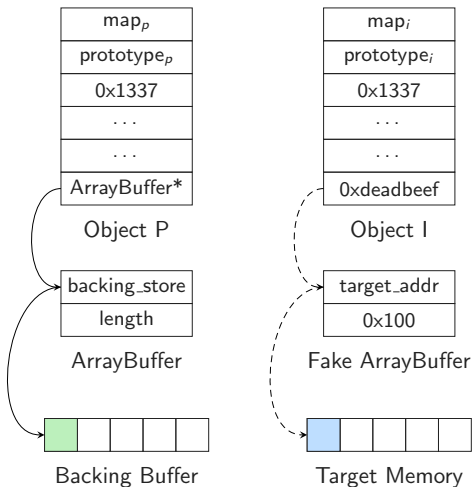
```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: 0xdeadbeef,  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: 0xdeadbeef,  
};
```

```
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



```
obj = {  
    i: 0x1337,  
    ...,  
    ptr: 0xdeadbeef,  
};  
  
if (obj instanceof ObjP) {  
    leak(obj.ptr[0])  
}
```



**DEMO**



- **Leak** 120 bit/hour over the **network**




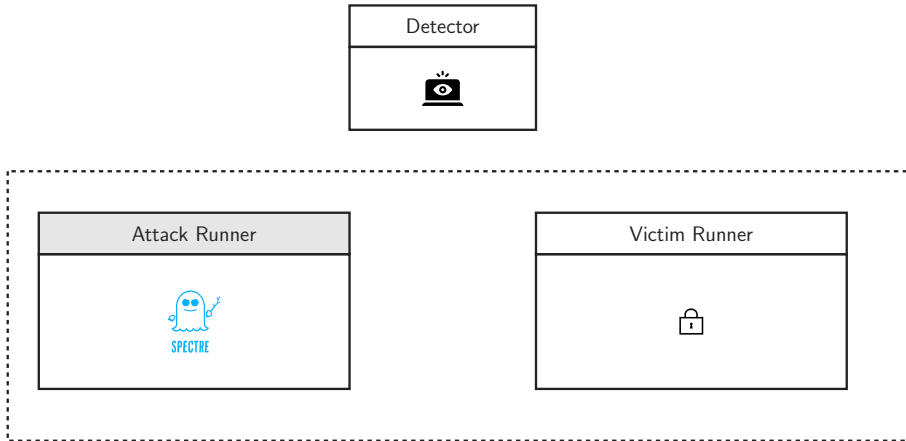
- **Leak** 120 bit/hour over the **network**
- Works within the offered runtime of 30 seconds

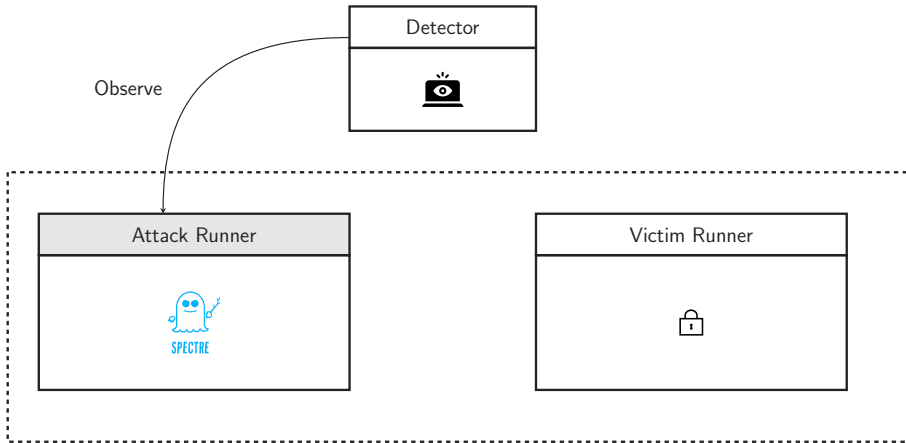


- **Leak** 120 bit/hour over the **network**
- Works within the offered runtime of 30 seconds
- Use speculative **type confusion** to create 64-bit leak primitive

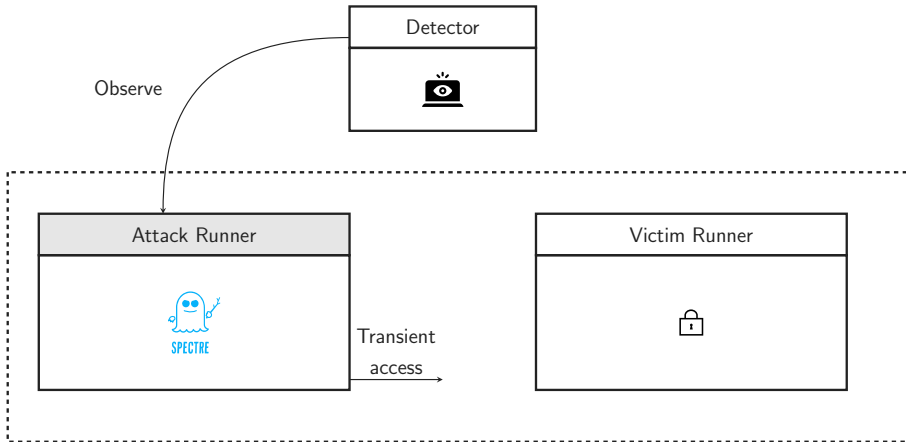
**Defense**

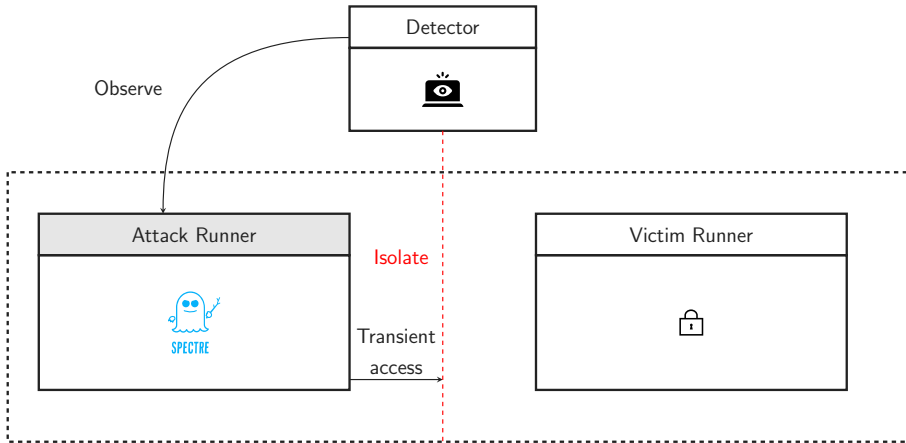


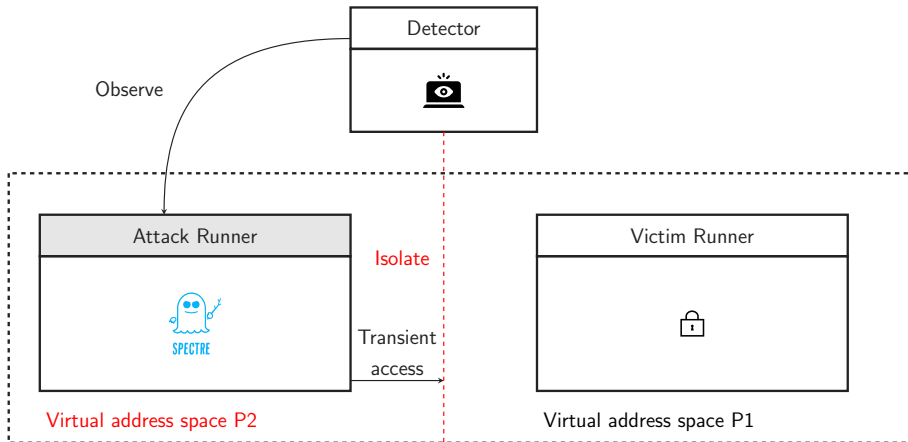














- How to **observe** the Attacker?


- How to **observe** the Attacker?
  - ✓ Performance Monitoring Counters



- How to **observe** the Attacker?
  - ✓ Performance Monitoring Counters
- Branch **Misses**



- How to **observe** the Attacker?
  - ✓ Performance Monitoring Counters
- Branch **Misses**
- Branch **Accesses**



- How to **observe** the Attacker?
  - ✓ Performance Monitoring Counters
- Branch **Misses**
- Branch **Accesses**
- Per **Code** Executed






- What is the **overhead**?




- What is the **overhead**?
- Different **Interfaces**
  - **PERF**
  - **rdmsr**
  - **rdpmc**




- What is the **overhead**?
  - Different **Interfaces**
    - **PERF**
    - **rdmsr**
    - **rdpmc**
- ✓ **rdpmc** → 2% overhead




- What is the **success rate**?




- What is the **success rate**?
- Detection vs leakage rate



- What is the **success rate**?
- Detection vs leakage rate
- ✓ Strong **reduction**




- Are there **false positives**?




- Are there **false positives**?
- Different CPUs







- Are there **false positives**?
- Different CPUs
- ✓ 4096 → 0.61%




- Can the observer be **tricked**?




- Can the observer be **tricked**?
- Additional Code Pages



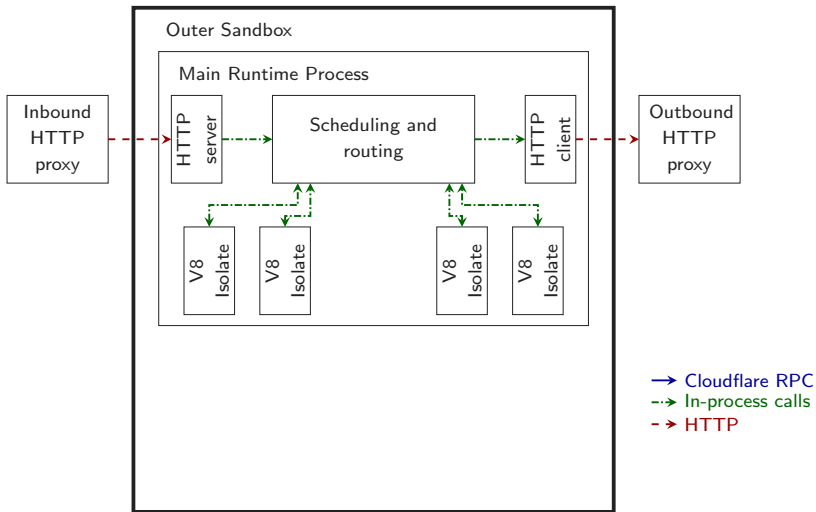
- Can the observer be **tricked**?
- Additional Code Pages
- Increase ITLB Accesses

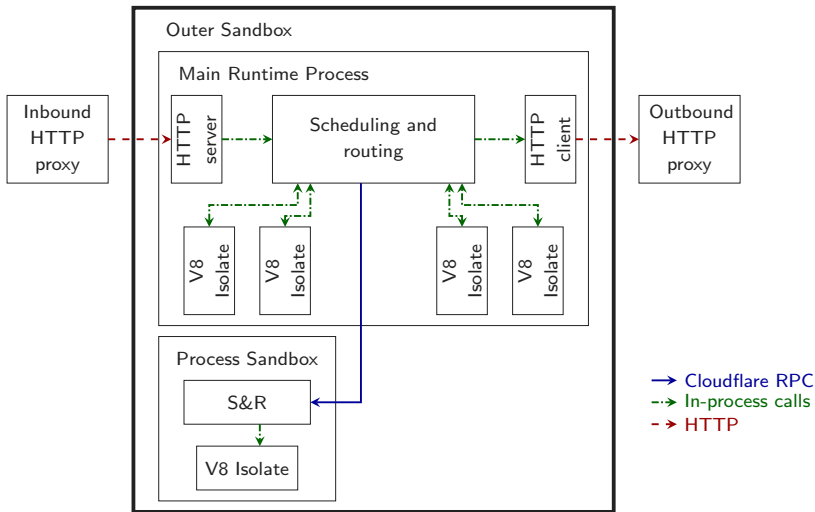


- Can the observer be **tricked**?
- Additional Code Pages
- Increase ITLB Accesses
- Reduction factor

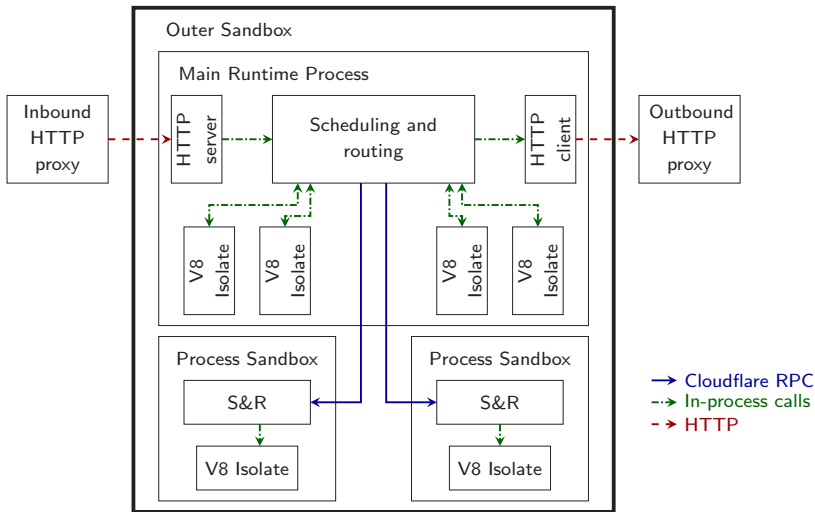


- Can the observer be **tricked**?
- Additional Code Pages
- Increase ITLB Accesses
- Reduction factor
- ✓ Not **compiled** by JavaScript









- Remote Spectre attacks **were** possible on Cloudflare Workers



- Remote Spectre attacks **were** possible on Cloudflare Workers
- Our solution detects all state-of-the-art Spectre attacks





- Remote Spectre attacks **were** possible on Cloudflare Workers
- Our solution detects all state-of-the-art Spectre attacks
- **Production-deployed** with a low false-positive rate



- Remote Spectre attacks **were** possible on Cloudflare Workers
- Our solution detects all state-of-the-art Spectre attacks
- **Production-deployed** with a low false-positive rate

## Dynamic Process Isolation

Martin Schwarzl

Graz University of Technology  
martin.schwarzl@iaik.tugraz.at

Pietro Borrello

Sapienza University of Rome  
borrello@diag.uniroma1.it

Andreas Kogler

Graz University of Technology  
andreas.kogler@iaik.tugraz.at

Kenton Varda

Cloudflare  
kenton@cloudflare.com

Thomas Schuster

Graz University of Technology  
thomas.schuster@student.tugraz.at

Daniel Gruss

Graz University of Technology  
daniel.gruss@iaik.tugraz.at

Michael Schwarz

CISPA Helmholtz Center for Information Security  
michael.schwarz@cispa.saarland